

```
!pip install xgboost
```

Requirement already satisfied: xgboost in c:\anaconda3\lib\site-packages (3.0.2)

Requirement already satisfied: numpy in c:\anaconda3\lib\site-packages (from xgboost) (1.26.4)

Requirement already satisfied: scipy in c:\anaconda3\lib\site-packages (from xgboost) (1.13.1)

```
!pip install shap
```

Requirement already satisfied: shap in c:\anaconda3\lib\site-packages (0.48.0)

Requirement already satisfied: numpy in c:\anaconda3\lib\site-packages (from shap) (1.26.4)

Requirement already satisfied: scipy in c:\anaconda3\lib\site-packages (from shap) (1.13.1)

Requirement already satisfied: scikit-learn in c:\anaconda3\lib\site-packages (from shap) (1.5.1)

Requirement already satisfied: pandas in c:\anaconda3\lib\site-packages (from shap) (2.2.2)

Requirement already satisfied: tqdm<=4.27.0 in c:\anaconda3\lib\site-packages (from shap) (4.66.5)

Requirement already satisfied: packaging>20.9 in c:\anaconda3\lib\site-packages (from shap) (24.1)

Requirement already satisfied: slicer==0.0.8 in c:\anaconda3\lib\site-packages (from shap) (0.0.8)

Requirement already satisfied: numba>=0.54 in c:\anaconda3\lib\site-packages (from shap) (0.60.0)

Requirement already satisfied: cloudpickle in c:\anaconda3\lib\site-packages (from shap) (3.0.0)

Requirement already satisfied: typing-extensions in c:\anaconda3\lib\site-packages (from shap) (4.11.0)

Requirement already satisfied: llvmlite<0.44,>=0.43.0dev0 in c:\anaconda3\lib\site-packages (from numba>=0.54->shap) (0.43.0)

Requirement already satisfied: colorama in c:\anaconda3\lib\site-packages (from tqdm<=4.27.0->shap) (0.4.6)

Requirement already satisfied: python-dateutil>=2.8.2 in c:\anaconda3\lib\site-packages (from pandas->shap) (2.9.0.post0)

Requirement already satisfied: pytz>=2020.1 in c:\anaconda3\lib\site-packages (from pandas->shap) (2024.1)

Requirement already satisfied: tzdata>=2022.7 in c:\anaconda3\lib\site-packages (from pandas->shap) (2023.3)

Requirement already satisfied: joblib>=1.2.0 in c:\anaconda3\lib\site-packages (from scikit-learn->shap) (1.4.2)

Requirement already satisfied: threadpoolctl>=3.1.0 in c:\anaconda3\lib\site-packages (from scikit-learn->shap) (3.5.0)

Requirement already satisfied: six>=1.5 in c:\anaconda3\lib\site-packages (from python-dateutil>=2.8.2->pandas->shap) (1.16.0)

```

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import LabelEncoder, OneHotEncoder,
StandardScaler
from sklearn.impute import SimpleImputer
from sklearn.model_selection import train_test_split, StratifiedKFold
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from xgboost import XGBClassifier
from sklearn.metrics import accuracy_score, f1_score, roc_auc_score,
confusion_matrix, classification_report
import shap
import warnings
warnings.filterwarnings("ignore")

crx_data_path = r"C:\Users\aniru\OneDrive\Desktop\ML tutorial\Credit
adjuster\crx.data"
crx_names_path = r"C:\Users\aniru\OneDrive\Desktop\ML tutorial\Credit
adjuster\crx.names"
index_path = r"C:\Users\aniru\OneDrive\Desktop\ML tutorial\Credit
adjuster\Index"
credit_lisp_path = r"C:\Users\aniru\OneDrive\Desktop\ML tutorial\
Credit adjuster\credit.lisp"
credit_names_path = r"C:\Users\aniru\OneDrive\Desktop\ML tutorial\
Credit adjuster\credit.names"

column_names = [
    'A1', 'A2', 'A3', 'A4', 'A5', 'A6', 'A7', 'A8',
    'A9', 'A10', 'A11', 'A12', 'A13', 'A14', 'A15', 'Class'
]
df_credit = pd.read_csv(crx_data_path, names=column_names,
na_values='?')

with open(crx_names_path, 'r', encoding='latin-1') as f:
    crx_names = f.read()
with open(index_path, 'r', encoding='latin-1') as f:
    index_file = f.read()
with open(credit_lisp_path, 'r', encoding='latin-1') as f:
    credit_lisp = f.read()
with open(credit_names_path, 'r', encoding='latin-1') as f:
    credit_names = f.read()

print("Files loaded successfully.")
print("DataFrame shape:", df_credit.shape)
print("crx.names preview:\n", crx_names[:300])

Files loaded successfully.
DataFrame shape: (690, 16)

```

crx.names preview:

1. Title: Credit Approval

2. Sources:

(confidential)

Submitted by quinlan@cs.su.oz.au

3. Past Usage:

See Quinlan,

\* "Simplifying decision trees", Int J Man-Machine Studies 27,  
Dec 1987, pp. 221-234.

\* "C4.5: Programs for Machine Learning", Morgan Kaufmann, Oct 1992

```
df_clean = df_credit.copy()
```

```
categorical_cols =
```

```
df_clean.select_dtypes(include='object').columns.tolist()
```

```
categorical_cols.remove('Class')
```

```
numerical_cols = df_clean.columns.difference(categorical_cols +  
['Class'])
```

```
cat_imputer = SimpleImputer(strategy='most_frequent')
```

```
num_imputer = SimpleImputer(strategy='mean')
```

```
df_clean[categorical_cols] =
```

```
cat_imputer.fit_transform(df_clean[categorical_cols])
```

```
df_clean[numerical_cols] =
```

```
num_imputer.fit_transform(df_clean[numerical_cols])
```

```
label_encoders = {}
```

```
for col in categorical_cols:
```

```
    le = LabelEncoder()
```

```
    df_clean[col] = le.fit_transform(df_clean[col])
```

```
    label_encoders[col] = le
```

```
df_clean['Class'] = df_clean['Class'].map({'+': 1, '-': 0})
```

```
scaler = StandardScaler()
```

```
df_clean[numerical_cols] =
```

```
scaler.fit_transform(df_clean[numerical_cols])
```

```
print("Data cleaned and preprocessed.")
```

```
print("Feature matrix shape:", df_clean.shape)
```

```
df_clean.head()
```

Data cleaned and preprocessed.

Feature matrix shape: (690, 16)

	A1	A2	A3	A4	A5	A6	A7	A8	A9	A10	A11
A12 \											
0	1	-0.062321	-0.956613	1	0	12	7	-0.291083	1	1	-0.288101
0											
1	0	2.288101	-0.060051	1	0	10	3	0.244190	1	1	0.740830
0											
2	0	-0.596738	-0.856102	1	0	10	3	-0.216324	1	0	-0.493887
0											
3	1	-0.315599	-0.647038	1	0	12	7	0.456505	1	1	0.535044
1											
4	1	-0.962303	0.174141	1	0	12	7	-0.153526	1	0	-0.493887
0											

	A13	A14	A15	Class
0	0	0.104544	-0.195413	1
1	0	-0.819689	-0.087852	1
2	0	0.557942	-0.037144	1
3	0	-0.488360	-0.194837	1
4	2	-0.372104	-0.195413	1

```
X = df_clean.drop('Class', axis=1)
y = df_clean['Class']
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42, stratify=y)
```

```
models = {
    "Logistic Regression": LogisticRegression(max_iter=1000),
    "Random Forest": RandomForestClassifier(n_estimators=100,
random_state=42),
    "XGBoost": XGBClassifier(use_label_encoder=False,
eval_metric='logloss', random_state=42)
}
```

```
for name, model in models.items():
    print(f"\nTraining: {name}")
    model.fit(X_train, y_train)
    preds = model.predict(X_test)
    probs = model.predict_proba(X_test)[: , 1]
    print("Accuracy:", round(accuracy_score(y_test, preds), 3))
    print("F1 Score:", round(f1_score(y_test, preds), 3))
    print("ROC AUC:", round(roc_auc_score(y_test, probs), 3))
    print("Classification Report:\n", classification_report(y_test,
preds))
```

```
Training: Logistic Regression
Accuracy: 0.899
F1 Score: 0.889
ROC AUC: 0.958
Classification Report:
```

	precision	recall	f1-score	support
0	0.93	0.88	0.91	77
1	0.86	0.92	0.89	61
accuracy			0.90	138
macro avg	0.90	0.90	0.90	138
weighted avg	0.90	0.90	0.90	138

Training: Random Forest

Accuracy: 0.891

F1 Score: 0.872

ROC AUC: 0.955

Classification Report:

	precision	recall	f1-score	support
0	0.88	0.94	0.91	77
1	0.91	0.84	0.87	61
accuracy			0.89	138
macro avg	0.89	0.89	0.89	138
weighted avg	0.89	0.89	0.89	138

Training: XGBoost

Accuracy: 0.906

F1 Score: 0.891

ROC AUC: 0.946


Classification Report:

	precision	recall	f1-score	support
0	0.90	0.94	0.92	77
1	0.91	0.87	0.89	61
accuracy			0.91	138
macro avg	0.91	0.90	0.90	138
weighted avg	0.91	0.91	0.91	138


```
explainer = shap.Explainer(models["XGBoost"])
```

```
shap_values = explainer(X_test)
```

```
shap.plots.beeswarm(shap_values)
```



```
shap.plots.waterfall(shap_values[0])
```



```

from sklearn.metrics import accuracy_score, f1_score, roc_auc_score,
classification_report

X = df_clean.drop('Class', axis=1)
y = df_clean['Class']

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42, stratify=y
)

models = {
    "Logistic Regression": LogisticRegression(max_iter=1000),
    "Random Forest": RandomForestClassifier(n_estimators=100,
random_state=42),
    "XGBoost": XGBClassifier(use_label_encoder=False,
eval_metric='logloss', random_state=42)
}

for name, model in models.items():
    print(f"\nTraining: {name}")
    model.fit(X_train, y_train)
    preds = model.predict(X_test)
    probs = model.predict_proba(X_test)[: , 1]
    print("Accuracy:", round(accuracy_score(y_test, preds), 3))
    print("F1 Score:", round(f1_score(y_test, preds), 3))
    print("ROC AUC:", round(roc_auc_score(y_test, probs), 3))
    print("Classification Report:\n", classification_report(y_test,
preds))

```

Training: Logistic Regression

Accuracy: 0.899

F1 Score: 0.889

ROC AUC: 0.958

Classification Report:

	precision	recall	f1-score	support
0	0.93	0.88	0.91	77
1	0.86	0.92	0.89	61
accuracy			0.90	138
macro avg	0.90	0.90	0.90	138
weighted avg	0.90	0.90	0.90	138

Training: Random Forest

Accuracy: 0.891

F1 Score: 0.872

ROC AUC: 0.955

Classification Report:

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

	0	0.88	0.94	0.91	77
	1	0.91	0.84	0.87	61
accuracy				0.89	138
macro avg		0.89	0.89	0.89	138
weighted avg		0.89	0.89	0.89	138

Training: XGBoost

Accuracy: 0.906

F1 Score: 0.891

ROC AUC: 0.946

Classification Report:

		precision	recall	f1-score	support
	0	0.90	0.94	0.92	77
	1	0.91	0.87	0.89	61
accuracy				0.91	138
macro avg		0.91	0.90	0.90	138
weighted avg		0.91	0.91	0.91	138

```
xgb_model = XGBClassifier(use_label_encoder=False,
eval_metric='logloss', random_state=42)
xgb_model.fit(X_train, y_train)
```

```
XGBClassifier(base_score=None, booster=None, callbacks=None,
               colsample_bylevel=None, colsample_bynode=None,
               colsample_bytree=None, device=None,
               early_stopping_rounds=None,
               enable_categorical=False, eval_metric='logloss',
               feature_types=None, feature_weights=None, gamma=None,
               grow_policy=None, importance_type=None,
               interaction_constraints=None, learning_rate=None,
               max_bin=None,
               max_cat_threshold=None, max_cat_to_onehot=None,
               max_delta_step=None, max_depth=None, max_leaves=None,
               min_child_weight=None, missing=nan,
               monotone_constraints=None,
               multi_strategy=None, n_estimators=None, n_jobs=None,
               num_parallel_tree=None, ...)
```

```
df_cross = df_clean.copy()
df_cross['approval_prob'] = xgb_model.predict_proba(X)[: , 1]
df_cross['decision'] = df_cross['approval_prob'].apply(
    lambda p: 'Approve' if p >= 0.6 else 'Manual Review'
)
```



```

def assign_risk(prob):
    if prob >= 0.8:
        return 'Low Risk'
    elif prob >= 0.6:
        return 'Moderate Risk'
    else:
        return 'High Risk'
df_cross['risk_bucket'] = df_cross['approval_prob'].apply(assign_risk)

df_cross['business_rule_flag'] = df_cross['risk_bucket'].apply(lambda
x: 1 if x == 'High Risk' else 0)

df_cross['final_decision'] = df_cross.apply(
    lambda row: 'Deny' if row['business_rule_flag'] == 1 else
row['decision'],
    axis=1
)

df_cross[['approval_prob', 'decision', 'risk_bucket',
'business_rule_flag', 'final_decision']].head()

```

	approval_prob	decision	risk_bucket	business_rule_flag
final_decision				
0	0.833003	Approve	Low Risk	0
Approve				
1	0.999512	Approve	Low Risk	0
Approve				
2	0.998048	Approve	Low Risk	0
Approve				
3	0.988194	Approve	Low Risk	0
Approve				
4	0.982082	Approve	Low Risk	0
Approve				

```


explainer = shap.Explainer(xgb_model)
shap_values = explainer(X)

shap.plots.beeswarm(shap_values)

```



```
sample_index = 0  
shap.plots.waterfall(shap_values[sample_index])
```



```
group_approval_rates = df_cross.groupby('A1')
['final_decision'].value_counts(normalize=True).unstack()
print(group_approval_rates)
```

final_decision	Approve	Deny
A1		
0	0.457143	0.542857
1	0.429167	0.570833

```
disparity = abs(group_approval_rates.loc[0, 'Approve'] -
group_approval_rates.loc[1, 'Approve'])
print(f"Disparate impact (A1): {round(disparity, 3)}")
```

Disparate impact (A1): 0.028

```
df_cross['needs_review'] = df_cross.apply(
    lambda row: 1 if (row['approval_prob'] < 0.6 or
row['business_rule_flag'] == 1) else 0,
    axis=1
)
```

```
df_cross[['approval_prob', 'final_decision', 'risk_bucket',
'business_rule_flag', 'needs_review']].head()
```

	approval_prob	final_decision	risk_bucket	business_rule_flag
needs_review				
0	0.833003	Approve	Low Risk	0
0				
1	0.999512	Approve	Low Risk	0
0				
2	0.998048	Approve	Low Risk	0
0				
3	0.988194	Approve	Low Risk	0
0				
4	0.982082	Approve	Low Risk	0
0				

```
df_temporal = df_cross.copy()
```

```
df_temporal['application_time'] = pd.date_range(start='2022-01-01',
periods=len(df_temporal), freq='H')
```

```
df_temporal['application_day'] =
df_temporal['application_time'].dt.date
df_temporal['daily_applicant_count'] =
df_temporal.groupby('application_day')
['application_day'].transform('count')
```

```
high_risk_times = df_temporal[df_temporal['risk_bucket'] == 'High
Risk']['application_time']
df_temporal['time_since_last_high_risk'] =
df_temporal['application_time'].apply(
    lambda t: (t - high_risk_times[high_risk_times <
t]).max()).total_seconds() / 3600
    if not high_risk_times[high_risk_times < t].empty else np.nan
)
```

```
df_temporal['rolling_approval_24h'] =
df_temporal['Class'].rolling(window=24, min_periods=1).mean()
```

```
df_temporal[['application_time', 'daily_applicant_count',
'time_since_last_high_risk', 'rolling_approval_24h']].head()
```

	application_time	daily_applicant_count
time_since_last_high_risk \		
0	2022-01-01 00:00:00	24
NaN		
1	2022-01-01 01:00:00	24
NaN		
2	2022-01-01 02:00:00	24
NaN		
3	2022-01-01 03:00:00	24
NaN		
4	2022-01-01 04:00:00	24
NaN		

	rolling_approval_24h
0	1.0
1	1.0
2	1.0
3	1.0
4	1.0

```
df_anomaly = df_temporal.copy()
```

```
exclude_cols = ['application_time', 'application_day', 'risk_bucket',
'final_decision']
```

```
X_anomaly = df_anomaly.drop(columns=exclude_cols, errors='ignore')
```

```
X_anomaly = X_anomaly.select_dtypes(include=[np.number])
```

```
X_anomaly_clean = X_anomaly.dropna()
```

```
from sklearn.ensemble import IsolationForest
```

```
from sklearn.impute import SimpleImputer
```

```
imputer = SimpleImputer(strategy='mean')
```

```
X_anomaly_clean = pd.DataFrame(imputer.fit_transform(X_anomaly),
columns=X_anomaly.columns)
```

```
iso_forest = IsolationForest(contamination=0.05, random_state=42)
```

```
df_anomaly.loc[X_anomaly_clean.index, 'anomaly_flag'] =
```

```
iso_forest.fit_predict(X_anomaly_clean)
```

```
df_anomaly['is_anomaly'] = df_anomaly['anomaly_flag'].map({-1: 1, 1:
0})
```

```
df_anomaly[['approval_prob', 'risk_bucket', 'is_anomaly']].head()
```

	approval_prob	risk_bucket	is_anomaly
0	0.833003	Low Risk	0
1	0.999512	Low Risk	0
2	0.998048	Low Risk	0
3	0.988194	Low Risk	0
4	0.982082	Low Risk	0

```
X_cluster =
```

```
df_anomaly.select_dtypes(include=[np.number]).drop(columns=['anomaly_f
lag', 'is_anomaly'], errors='ignore')
```

```
X_cluster = X_cluster.fillna(X_cluster.mean())
```

```
df_cluster = df_anomaly.copy()
```

```
X_cluster = df_cluster.select_dtypes(include=[np.number])
```

```

X_cluster = X_cluster.dropna()
df_cluster = df_cluster.loc[X_cluster.index]

from sklearn.cluster import KMeans

kmeans = KMeans(n_clusters=3, random_state=42)
df_cluster['cluster_id'] = kmeans.fit_predict(X_cluster)

cluster_risk_summary = df_cluster.groupby('cluster_id')
['risk_bucket'].value_counts(normalize=True).unstack()

print(cluster_risk_summary)
df_cluster[['cluster_id', 'risk_bucket', 'is_anomaly']].head()

risk_bucket  High Risk  Low Risk  Moderate Risk
cluster_id
0             0.831486  0.164080             0.004435
1             0.028571  0.952381             0.019048
2             0.118421  0.881579             NaN

   cluster_id risk_bucket  is_anomaly
58           0     Low Risk           1
59           0     Low Risk           0
60           0     Low Risk           0
61           0     Low Risk           0
62           1     Low Risk           0

X =
df_cluster.select_dtypes(include=[np.number]).drop(columns=['Class'],
errors='ignore').dropna()
y = df_cluster.loc[X.index, 'Class']

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42, stratify=y
)

xgb_cost_sensitive = XGBClassifier(
    use_label_encoder=False,
    eval_metric='logloss',
    random_state=42,
    scale_pos_weight=2.0 # Adjust based on business risk profile
)

xgb_cost_sensitive.fit(X_train, y_train)

XGBClassifier(base_score=None, booster=None, callbacks=None,
               colsample_bylevel=None, colsample_bynode=None,
               colsample_bytree=None, device=None,
early_stopping_rounds=None,
               enable_categorical=False, eval_metric='logloss',
               feature_types=None, feature_weights=None, gamma=None,
               grow_policy=None, importance_type=None,

```

```

max_bin=None, interaction_constraints=None, learning_rate=None,
max_cat_threshold=None, max_cat_to_onehot=None,
max_delta_step=None, max_depth=None, max_leaves=None,
min_child_weight=None, missing=nan,
monotone_constraints=None,
multi_strategy=None, n_estimators=None, n_jobs=None,
num_parallel_tree=None, ...)

```

```

y_pred = xgb_cost_sensitive.predict(X_test)
y_prob = xgb_cost_sensitive.predict_proba(X_test)[:, 1]
print("Accuracy:", round(accuracy_score(y_test, y_pred), 3))
print("F1 Score:", round(f1_score(y_test, y_pred), 3))
print("ROC AUC:", round(roc_auc_score(y_test, y_prob), 3))
print("Classification Report:\n", classification_report(y_test,
y_pred))

```

Accuracy: 0.921

F1 Score: 0.909

ROC AUC: 0.995

Classification Report:

	precision	recall	f1-score	support
0	1.00	0.87	0.93	77
1	0.83	1.00	0.91	50
accuracy			0.92	127
macro avg	0.92	0.94	0.92	127
weighted avg	0.93	0.92	0.92	127

```
df_experimental = df_cluster.copy()
```

```
df_experimental['credit_utilization'] = df_experimental['A14'] /
(df_experimental['A15'] + 1)
```

```
df_experimental['is_new_customer'] = (df_experimental['A11'] ==
0).astype(int)
```

```
df_experimental['has_missed_payments'] = (df_experimental['A12'] ==
1).astype(int)
```

```
df_experimental[['A14', 'A15', 'credit_utilization', 'A11',
'is_new_customer', 'A12', 'has_missed_payments']].head()
```

	A14	A15	credit_utilization	A11	is_new_customer
A12 \					
58	3.063254	-0.195413	3.807239	-0.493887	0
1					
59	-0.604616	-0.195413	-0.751461	1.769760	0
0					

60	0.092919	-0.195413	0.115486	0.329258	0
1					
61	0.383558	-0.055199	0.405967	1.358188	0
0					
62	-0.604616	-0.118584	-0.685959	-0.082314	0
1					

	has_missed_payments
58	1
59	0
60	1
61	0
62	1

```
df_experimental['A12'] = df_experimental['A12'].map({'t': 1, 'f': 0})
```

```
df = df_experimental.copy()
X = df.select_dtypes(include=[np.number]).drop(columns=['Class'],
errors='ignore').dropna()
y = df.loc[X.index, 'Class']
df = df.loc[X.index]
```

```
print(y.value_counts(dropna=False))
```

```
Series([], Name: count, dtype: int64)
```

```
df['Class'] = df['Class'].map({'+': 1, '-': 0})
```

```
df = df[df['Class'].notna()]
```

```
assert set(y.unique()).issubset({0, 1}), "Your target is not binary"
```

```
print("X shape:", X.shape)
print("df['Class'] shape:", df['Class'].shape)
print("Any NaNs in Class?", df['Class'].isna().sum())
print("Unique Class values:", df['Class'].unique())
```

```
X shape: (0, 27)
df['Class'] shape: (0,)
Any NaNs in Class? 0
Unique Class values: []
```

```
df['Class'] = df['Class'].map({'+': 1, '-': 0})
df = df[df['Class'].notna()]
```

```
X = df.select_dtypes(include=[np.number]).drop(columns=['Class'],
errors='ignore').dropna()
y = df.loc[X.index, 'Class']
```

```
print("Target Distribution:")
print(y.value_counts())
```



```
Series([], Name: count, dtype: int64)
```

```
print(df['Class'].value_counts(dropna=False))
```

```
print(df['Class'].unique())
```

```
print(df['Class'].head())
```

```
Series([], Name: count, dtype: int64)
```

[ ]

```
Series([], Name: Class, dtype: int64)
```

```
print(df_experimental.columns)
```

```
Index(['A1', 'A2', 'A3', 'A4', 'A5', 'A6', 'A7', 'A8', 'A9', 'A10',
      'A11',
```

```
'A12', 'A13', 'A14', 'A15', 'Class', 'approval_prob',
```

```
'decision',
```

```
'risk_bucket', 'business_rule_flag', 'final_decision',
```

```
'needs_review',
```

```
    'application_time', 'application_day', 'daily_applicant_count',
```

```
'time_since_last_high_risk', 'rolling_approval_24h',
```

```
'anomaly_flag',
```

```
'is_anomaly', 'cluster_id', 'credit_utilization',
```

```
'is_new_customer',
```

```
    'has_missed_payments'],
```

```
dtype='object')
```

```
print(df_experimental['Class'].value_counts(dropna=False))
```

```
print(df_experimental['Class'].unique())
```

```
print(df_experimental['Class'].dtype)
```

## Class

0 383

1 249

```
Name: count, dtype: int64
```

 $[1 \ 0]$ 

```
int64
```

```
print("X shape:", X.shape)
```

```
print("Any NaNs in X?", X.isna().sum().sum())
```

```
print("X summary stats:\n", X.describe())
```

X shape: (0, 27)

Any NaNs in X? 0

X summary stats:

A1    A2    A3    A4    A5    A6    A7    A8    A9    A10    ...

```
needs review \
```

```
count= 0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  ...
```

0.0

```
mean      NaN     NaN     NaN     NaN     NaN     NaN     NaN     NaN     NaN     NaN     ...
```

NaN

std	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	...
NaN											
min	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	...
NaN											
25%	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	...
NaN											
50%	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	...
NaN											
75%	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	...
NaN											
max	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	...
NaN											

daily_applicant_count		time_since_last_high_risk	
rolling_approval_24h \			
count		0.0	0.0
0.0			
mean		NaN	NaN
NaN			
std		NaN	NaN
NaN			
min		NaN	NaN
NaN			
25%		NaN	NaN
NaN			
50%		NaN	NaN
NaN			
75%		NaN	NaN
NaN			
max		NaN	NaN
NaN			

anomaly_flag		is_anomaly		cluster_id		credit_utilization \	
count	0.0	0.0	0.0	0.0	0.0	0.0	
mean	NaN	NaN	NaN	NaN	NaN	NaN	
std	NaN	NaN	NaN	NaN	NaN	NaN	
min	NaN	NaN	NaN	NaN	NaN	NaN	
25%	NaN	NaN	NaN	NaN	NaN	NaN	
50%	NaN	NaN	NaN	NaN	NaN	NaN	
75%	NaN	NaN	NaN	NaN	NaN	NaN	
max	NaN	NaN	NaN	NaN	NaN	NaN	

is_new_customer		has_missed_payments	
count	0.0	0.0	0.0
mean	NaN	NaN	NaN
std	NaN	NaN	NaN
min	NaN	NaN	NaN
25%	NaN	NaN	NaN
50%	NaN	NaN	NaN
75%	NaN	NaN	NaN

```

max                NaN                NaN

[8 rows x 27 columns]

X =
df_experimental.select_dtypes(include=[np.number]).drop(columns=['Class'], errors='ignore')
X = X.fillna(X.mean())

y = df_experimental['Class'].astype(int)

print(X.shape)
print(y.value_counts())

(632, 27)
Class
0      383
1      249
Name: count, dtype: int64

model = XGBClassifier(use_label_encoder=False, eval_metric='logloss',
                      random_state=42)
model.fit(X, y)

XGBClassifier(base_score=None, booster=None, callbacks=None,
              colsample_bylevel=None, colsample_bynode=None,
              colsample_bytree=None, device=None,
              early_stopping_rounds=None,
              enable_categorical=False, eval_metric='logloss',
              feature_types=None, feature_weights=None, gamma=None,
              grow_policy=None, importance_type=None,
              interaction_constraints=None, learning_rate=None,
              max_bin=None,
              max_cat_threshold=None, max_cat_to_onehot=None,
              max_delta_step=None, max_depth=None, max_leaves=None,
              min_child_weight=None, missing=nan,
              monotone_constraints=None,
              multi_strategy=None, n_estimators=None, n_jobs=None,
              num_parallel_tree=None, ...)

df['approval_prob'] = model.predict_proba(X)[:, 1]
df['prediction'] = model.predict(X)
df['final_decision'] = df['approval_prob'].apply(lambda p: 'Approve'
if p >= 0.6 else 'Manual Review')
df['bias_flag'] = (df['A1'] == 0).astype(int)

explainer = shap.TreeExplainer(model)
shap_vals = explainer.shap_values(X)

top_shap_features = []
for row_vals in shap_vals:

```

```

top_idx = np.argsort(np.abs(row_vals))[:, :-1][:3]
top_feats = [X.columns[i] for i in top_idx]
top_shap_features.append(', '.join(top_feats))
df['top_3_shap_features'] = top_shap_features

audit_trail = df[[
    'prediction', 'approval_prob', 'final_decision',
    'risk_bucket', 'bias_flag', 'top_3_shap_features'
]]

audit_trail.head()

```

	prediction	approval_prob	final_decision	risk_bucket	bias_flag	\
0	1	0.988367	Approve	NaN	0	
1	1	0.998461	Approve	NaN	0	
2	1	0.997990	Approve	NaN	0	
3	1	0.999494	Approve	NaN	0	
4	1	0.994939	Approve	NaN	0	

	top_3_shap_features
0	approval_prob, rolling_approval_24h, time_sinc...
1	approval_prob, rolling_approval_24h, time_sinc...
2	approval_prob, rolling_approval_24h, time_sinc...
3	approval_prob, rolling_approval_24h, time_sinc...
4	approval_prob, rolling_approval_24h, time_sinc...